

Aryan Miriyala

+1 419-315-0444 | aryanmiriyala@gmail.com | [linkedin.com/in/aryan-miriyala](https://www.linkedin.com/in/aryan-miriyala) | github.com/aryanmiriyala

August 26, 2026

Hiring Team
Flash AI

Dear Hiring Team,

I am applying for the Full Stack Engineer role at Flash AI. The work caught my attention because it combines secure software, applied AI, and evidence-heavy search systems for public safety users. That is close to the kind of engineering I have been drawn toward: systems where grounding, auditability, access control, and reliability matter as much as the model output.

At Actual Reality Technologies, I built AI Gateway routing and safety infrastructure around LiteLLM, including smart-router tiers, guardrails, Langfuse tracing, Slack block alerts, CI gates, and red-team fixtures. That work pushed me to think about LLM calls as production systems: what should be admitted, how routing decisions should be tested, how failures should be visible, and how safety behavior can be reviewed instead of treated as a black box.

I also have the full-stack and data background that maps to your stack. I built FalconGraph, a source-grounded RAG search system using Python, FastAPI, FAISS retrieval, document/web crawling, and a Next.js interface with cited answers for natural language questions. In internships, I have worked with PostgreSQL, AWS Glue, S3, IAM, Lambda, Docker-based development workflows, Angular, React, Node.js, secure authentication, RBAC/JWT, anti-CSRF protection, audit logging, and large production data pipelines.

I would be excited to bring that mix of Python/backend engineering, RAG/search work, frontend product experience, and security-aware development to Flash AI as you build tools for investigators and analysts.

Sincerely,

Aryan Miriyala