Aryan Miriyala

aryanmiriyala@gmail.com | +1 419-315-0444 | linkedin.com/in/aryan-miriyala | github.com/aryanmiriyala

Dear Pulse Hiring Team,

Pulse's focus on document understanding at scale — combining layout detection, low-latency OCR, reading-order algorithms, table parsing, and fine-tuned vision-language models — is exactly the kind of applied AI work I want to be doing. Most parsing tools treat documents as flat text; you are building a multi-stage pipeline that actually understands structure. That distinction matters because it turns extraction into a reliable engineering problem rather than a prompting exercise.

My background aligns with that direction in three ways. At Actual Reality Technologies, I built and operate the AI Gateway infrastructure around a LiteLLM proxy: multi-tier model routing, guardrails, Langfuse tracing, Slack block alerts, CI/CD gates, and 44 red-team fixtures. I also built the automated evaluation pipeline that validates the IRIS glass-research AI chatbot across 41+ regression test suites using Node.js and Playwright, verifying multi-turn reasoning, tool execution, and citation accuracy against thermodynamic calculations and ASTM standards. This maps directly to Pulse's need for production-grade extraction pipelines with measurable accuracy and latency guarantees.

On the full-stack side, I architected and delivered the production customer portal end-to-end using Next.js 15, React 19, TypeScript, and Firebase — implementing multi-tenant session isolation, two-way Plane API ticket reconciliation with cursor-based sync, and real-time client retainer tracking. Earlier at SmartSolve, I built a secure onboarding system with Next.js, TypeScript, PostgreSQL, Drizzle ORM, SSO, and RBAC middleware. At AAIS, I engineered automated ETL pipelines with Python, PySpark, and AWS Glue processing 20+ TB across 700+ member companies, and built serverless validation workflows with AWS Lambda and S3.

For document intelligence specifically, my FalconGraph Search project built a source-grounded RAG system that crawls web resources, processes PDF/DOCX/HTML documents through a FastAPI pipeline, generates embeddings, retrieves with FAISS, and serves cited answers through a Next.js interface. The pipeline handles document ingestion, chunking, and semantic ranking — adjacent work to the layout understanding and table parsing challenges you are solving.

I am based in Ohio and ready to relocate to San Francisco for the on-site role. I would welcome the chance to discuss how my experience building AI gateway infrastructure, evaluation pipelines, and full-stack production systems can contribute to Pulse's extraction platform.

Sincerely,

Aryan Miriyala